

Codebook

The IUROPA CJEU Database Platform

Joshua C. Fjelstul, Ph.D.

March 31, 2026

About IUROPA

IUROPA (*ius + Europa*) is a multidisciplinary platform for research on judicial politics in the European Union (EU). It brings together a network of scholars with backgrounds in law and political science who share an interest in research on law and politics related to the Court of Justice of the European Union (CJEU).

A key feature of IUROPA is the development of a comprehensive, reliable, and up-to-date database with detailed information on cases, decisions, and actors involved in the judicial processes of the CJEU: the IUROPA CJEU Database Platform.

About the Database

The **IUROPA CJEU Database Platform** is the largest and most comprehensive research-oriented database on the CJEU. It includes data on the universe of CJEU cases, proceedings, decisions, judges, and more, collected from the Registry of the Court, InfoCuria (the Court's official database), and EUR-Lex (the EU's official database of legal documents). Some of the data is hand-coded, and some is machine-coded from metadata or the official French-language text of decisions. All of the data has been cleaned and cross-validated and is research-ready.

The database addresses three fundamental problems with data on CJEU decisions from official sources that make quantitative analysis difficult. First, there are inconsistencies across official data sources with respect to basic information like decision IDs, case IDs, proceeding IDs, and decision dates. Second, ECLI numbers do not always uniquely identify decisions. Third, it is difficult to accurately identify exactly which cases and proceedings each decision applies to, because of how the Court joins cases together.

License

The IUROPA CJEU Database Platform is published under a **CC-BY-NC-SA 4.0 license** (Attribution-NonCommercial-ShareAlike). You can use and share the data for non-commercial purposes as long as you provide proper attribution, provide a link to the license, and indicate if you made any changes to the database.

To provide proper attribution under the CC-BY-NC-SA 4.0 license, you must provide the names of the authors, a copyright notice, a license notice, a disclaimer notice, and a link to the database. If you remix, transform, or build upon the database, you must distribute your contribution under the same license as the original.

Consistent with the CC-BY-NC-SA 4.0 license, we provide this database as-is and as-available, and make no representations or warranties of any kind concerning the database, whether express, implied, statutory, or other. This includes, without limitation, warranties of title, merchantability, fitness for a particular purpose, non-infringement, absence of latent or other defects, accuracy, or the presence or absence of errors, whether or not known or discoverable.

Database Structure

Version 3.0 of the IUROPA CJEU Database Platform is a relational database. Each table has a distinct unit of observation. Some tables have a **simple unit of observation**, where each row is uniquely identified by a single primary key column. Others have a **complex unit of observation**, where each row is uniquely identified by a combination of two or more columns (a composite primary key). One or more columns in a composite key will typically be a foreign key — a column that serves as the primary key in another table.

For example, in the `decisions` table, which has one observation per unique decision, `decision_id` serves as the primary key. In the `judge_assignments` table, which has one observation per judge per decision, each observation is uniquely defined by the combination of `decision_id` and `judge_id`. Both of these are foreign keys: `decision_id` references `decisions` and `judge_id` references `judges_advocates_general`.

Cases and Proceedings

The database makes an important distinction between **cases** and **proceedings**.

A **case** corresponds to a case number assigned by the Court when a case is lodged (e.g., `C-370/12`). Case numbers are recorded in `case_number` and follow the format `X-####/##`, where `X` indicates the court, the first number is the case number within the year, and the second number is the last two digits of the year.

A **proceeding** is an action before the Court involving a specific set of cases and legal procedures. A case always has a main proceeding (e.g., an action for annulment, a reference for a preliminary ruling), and that proceeding can include multiple legal procedures. A case can also have secondary proceedings — for example, an application for interim measures, an application as to costs, or an application for legal aid. Secondary proceedings always have a suffix in their proceeding ID.

The Court can also **join proceedings** together, meaning a proceeding may involve multiple cases. When cases are joined, the proceeding number uses the case number of the case that was lodged first (the lead case). Proceeding numbers can also have suffixes that indicate the

legal procedure involved (e.g., `P` for appeals, from the French *pourvoi*; `DEP` for applications as to costs). See the `proceeding_suffixes` table for the full list of suffixes.

This case/proceeding distinction matters for merging. The parties in secondary proceedings can differ from the parties in the main proceeding. Analysis at the case level, proceeding level, and decision level can all yield different results, and the right level depends on the research question.

ECLI Numbers and Decision IDs

There are two general problems with ECLI numbers in the official sources: (1) sometimes one ECLI number is assigned to more than one decision, and (2) sometimes more than one ECLI number is assigned to one decision. These problems particularly affect orders. By cross-referencing data from the Registry, InfoCuria, and EUR-Lex, the database addresses both problems and constructs a complete table of unique decisions.

Each decision is assigned a `decision_id` (the IUROPA decision ID) that encodes the court, the case of record, the proceeding suffix (if applicable), the decision type, and the date. This makes `decision_id` more informative than an ECLI number, which encodes none of this information and cannot be reconstructed from decision metadata alone.

Unique Identifiers

The database uses a consistent set of string-based ID columns throughout. These are more informative than numeric surrogate keys because they encode meaningful information (court, year, case number, decision type). The key IDs are: `case_id`, which uniquely identifies each case; `proceeding_id`, which uniquely identifies each proceeding; `decision_id`, which uniquely identifies each decision; and `actor_id`, which uniquely identifies each actor, including judges and AGs.

Merging Tables

The IUROPA CJEU Database Platform is a relation database. The database is normalized, meaning that information is not repeated across tables. To analyze, say, the formation of the panel that decided a judgment alongside the outcome of the case, you will need to join `decisions`, `judge_assignments`, `chamber_assignments`, and `procedures`. The curated tables handle many of these joins for common use cases.

When joining tables, always use one of the IUROPA ID numbers, like `decision_id`, `proceeding_id`, `case_id`, or `actor_id`. ECLI or CELEX numbers are not always uniquely

identifying, and using them for merging can cause duplicate observations and/or dropped observations.

One-to-Many Relationships

Several tables have multiple observations per decision (e.g., `procedures`, `parties`, `judge_assignments`). If you join two such tables directly to `decisions`, you will get a Cartesian product — one row for every combination of, say, party and judge. If you want your table to still have one row per decision after the merge, you will need to aggregate these tables before merging, or merge them sequentially with care. The curated tables are designed to avoid this problem.

Joined Proceedings

When the Court joins proceedings together, decisions in the `decisions` table reference the lead proceeding (the earliest proceeding, or the proceeding with the lowest case number). The `joins` table records all joining relationships attested in the Registry or in the text of a judgment. If your research treats joined cases as separate units, you will need to explicitly expand using the `joins` table. If you are merging on `proceeding_id` and expecting one row per originally-filed case, you will miss cases that were joined to another proceeding.

Cases

One observation per unique case. A case number is a proceeding number without the suffix (e.g., C-1/26 vs C-1/26 DEP). If the proceeding number does not have a suffix, the case number and the proceeding number are the same.

Notes

- The purposes of the `cases` table is to show which proceedings are grouped together in the same case.
- The table includes every case number that is attested in the Registry, InfoCuria, or EUR-Lex. Not all cases are attested in all official sources. Some case numbers that should exist, based on the highest case number in each year, do not appear in any official source. Those are omitted.
- Do not merge the `cases` table directly with the `decisions` table. Decisions are linked to proceedings, not directly to cases. To connect cases to decisions, join `cases` to `proceedings` to `decisions`. If you try to merge `cases` and `decisions` together using `case_id` and `proceeding_id`, you will drop all observations where the `proceeding_id` includes a suffix, which includes all appeals and some references for a preliminary ruling.
- The `case_number` and `case_id` columns in the `cases` table omit all proceeding suffixes. Main proceedings with suffixes (e.g., `PPU`, `P`) and secondary proceedings with suffixes (e.g., `R`, `DEP`) are all associated with the same `case_id`. If you need to distinguish between main and secondary proceedings, work at the `proceedings` level.
- Court of Justice cases before the creation of the General Court do not use a `C-` prefix in official sources, but the database always includes the prefix for consistency.

Variables

- `case_id` `text` `primary key` `required`
A unique identifier for the case. Format: `CASE_X_####_####`, where `X` indicates the court (`CJ` for Court of Justice, `GC` for General Court, `CST` for Civil Service Tribunal), the first number is the case number within the year (with leading zeros), and the second number is the year.

- `case_number` `text` `required`

The case number assigned by the Court. Format: `X-####/##`, where `X` is a letter indicating the court (`C` for Court of Justice, `T` for General Court, `F` for Civil Service Tribunal), the first number is the case number within the year, and the second number is the last two digits of the year. For Court of Justice cases lodged before the creation of the General Court, the Court does not use a `C-` prefix, but the database always includes it for clarity. There are no missing values.

- `court` `text` `required`

The court at which the case was lodged. Possible values: `Court of Justice`, `General Court`, and `Civil Service Tribunal`.

- `prefix` `text` `required`

The prefix from the case number. Cases at the Court of Justice before the General Court was created do not have prefixes, but they have been added to all case numbers for consistency. Possible values: `C-`, `T-`, and `F-`.

- `year` `integer` `required`

The year the case was lodged.

Proceedings

One observation per unique proceeding. The term 'proceeding' has two meanings: (1) one of the proceedings within a case (e.g., an application for interim measures in an annulment case), or (2) a bundle of joined proceedings where the lowest case number becomes the lead case number for the bundle. Many proceedings have suffixes (e.g., R, P, DEP) indicating the legal procedure.

Notes

- Suffixes are used for secondary proceedings (e.g., applications for interim measures, applications as to costs) but are also used in some main proceedings (e.g., **PPU** for preliminary rulings that use the urgent procedure, or **P** for appeals).
- Proceeding names come from the Registry and are not available for all proceedings due to gaps in the official record.
- Proceeding names do not include all parties. The names of the parties that are included are often reduced (e.g., company suffixes dropped). For preliminary rulings, they may only include the name of the applicant at the national court.
- Proceeding names and case names are two different things. If a case has only one proceeding, the proceeding name can be used as a case name. But if there are multiple proceedings, it is possible for the parties to vary across proceedings. The case name is the proceeding name for the main proceeding.

Variables

- **proceeding_id** **text** **primary key** **required**
A unique identifier for the proceeding. Format: **PROC_X_####_####_Y**, where **X** indicates the court (**CJ** for Court of Justice, **GC** for General Court, **CST** for Civil Service Tribunal), the first number is the case number within the year (with leading 0s), the second number is the year, and **Y** is the suffix (normalized with underscores).
- **case_id** **text** **required**
The case this proceeding belongs to.

- `proceeding_number` `text` `required`

The proceeding number assigned by the Court. Format: `X-####/## Y`, where `X` is a letter indicating the court (`C` for Court of Justice, `T` for General Court, `F` for Civil Service Tribunal), the first number is the case number within the year, the second number is the last two digits of the year, and `Y` is the suffix. For Court of Justice cases lodged before the creation of the General Court, the Court does not use a `C-` prefix, but the database always includes it for clarity.

- `court` `text` `required`

The court at which the proceeding was lodged. Possible values: `Court of Justice`, `General Court`, and `Civil Service Tribunal`.

- `prefix` `text` `required`

The prefix from the proceeding number. Proceedings at the Court of Justice before the General Court was created do not have prefixes, but they have been added to all proceeding numbers for consistency. Possible values: `C-`, `T-`, and `F-`.

- `year` `integer` `required`

The year the proceeding was lodged.

- `suffix` `text` `nullable`

The proceeding suffix (e.g., `P` for appeals, `DEP` for costs, `R` for interim measures). Missing when there is no proceeding suffix.

- `proceeding_name` `text` `nullable`

The name of the proceeding as it appears in the Registry. Missing when the proceeding does not appear the official registry. This sometimes happens for proceedings that are joined to other proceedings.

Joins

One observation per joining of a proceeding to a lead proceeding. Records cases that have been joined for a single decision.

Notes

- Proceedings are joined, not cases. Cases are implicitly joined by virtue of one or more of their proceedings being joined. As such, joins are fundamentally a proceeding-level event. When a case has multiple proceedings, it is possible for only one of the proceedings to be joined. For example, secondary proceedings (e.g., applications for interim measures) in a set of cases can be joined even if the main proceedings are not.
- Every proceeding that involves joined proceedings has a lead proceeding number. This is the proceeding number that is used for the bundle of joined cases. It is always the case with the lowest number (e.g., the first of the cases to be lodged). The joined proceedings are the proceedings that are joined to the lead case.
- Data on joins comes from several sources. First, the Registry includes entries for joins. These sometimes include the ECLI number for the order that joined the cases. Second, the Registry, EUR-Lex metadata, and the text of judgments and orders include lists of all cases associated with decisions. This information is not always correct or complete. This table includes all joins that are known either from the Registry entries or from these lists of cases associated with each judgment or order.

Variables

- `joined_proceeding_id` `text` `required`
The proceeding that was joined to the lead proceeding.
- `lead_proceeding_id` `text` `required`
The lead proceeding to which the other proceeding was joined.

Transfers

One observation per transfer of a proceeding to a different court.

Notes

- When a proceeding is transferred from one court to another, it receives an entirely new proceeding number. It is not possible to tell whether two cases at two different courts are the same case based on the proceeding numbers.
- There are several situations in which proceedings have been transferred from one court to another. First, when the General Court was created (then called the Court of First Instance), some Court of Justice proceedings were transferred to it. Second, when the Civil Service Tribunal was created, some proceedings were transferred to it. Third, when the Civil Service Tribunal was disbanded, some proceedings were transferred back to the General Court or the Court of Justice.
- Data on transfers comes from the Registry, which records the original proceeding number and the new proceeding number.

Variables

- `original_proceeding_id` `text` `required`
The proceeding before the transfer.
- `new_proceeding_id` `text` `required`
The proceeding after the transfer, at the new court.

Appeals

One observation per proceeding in which a decision is appealed.

Notes

- Every appeal has a `P` suffix. There are a number of variants. `P` is an appeal of a proceeding. `P(I)` is an appeal against a decision in a proceeding concerning an application to intervene. `P(AJ)` is an appeal against a decision in a proceeding concerning an application for legal aid. `P(R)` is an appeal against a decision in a proceeding concerning an application for interim measures. There are also suffixes like `P-R`, which indicates an application for interim measures in an appeal, and `P-REV`, which indicates an application for the revision of a decision in an appeal. There are a number of other variants. See the `proceeding_suffixes` table for a complete list.
- There is a many-to-many relationship between original proceedings and appeal proceedings. One proceeding can generate one appeal, one proceeding can generate multiple appeals, and multiple proceedings can generate a single appeal.
- Every appeal should be associated with at least one original proceeding. There are a small number of situations where the original proceeding associated with an appeal is not available from official sources. Proceedings numbers for proceedings that are appealed are not marked in any way, so there is not a systematic way to look for missing appeals, if there are any.

Variables

- `original_proceeding_id` `text` `nullable`
The original proceeding being appealed. Missing when official sources, including the text of the decision, do not report the number of the original proceeding.
- `appeal_proceeding_id` `text` `required`
The appeal proceeding.

Proceeding suffixes

One observation per unique proceeding suffix.

Notes

- CJEU proceedings sometimes have suffixes (e.g., `P`, `DEP`) that tell you something about the legal procedure. This table explains what each suffix means. All of these suffixes have been standardized across the database to match this table (there is some variance in formatting in older proceedings). Proceeding suffixes can tell you about the main proceeding (e.g., `P` for appeals), secondary proceedings (e.g., `DEP` for applications as to costs), and about special procedures (e.g., `PPU` for the urgent preliminary ruling procedure).

Variables

- `proceeding_suffix` `text` `required`
The proceeding suffix.
- `description` `text` `required`
A description of what the suffix means.

Decisions

One observation per decision. Decisions include judgments, orders, and Advocate General (AG) opinions.

Notes

- The decisions table is the central table in the database. In particular, it links many other tables (e.g., `judge_assignments`, `procedures`) to the proceedings table via `proceeding_id`.
- The `decisions` table covers judgments, orders, and AG opinions. It covers the Court of Justice, the General Court, and the Civil Service Tribunal.
- The `decision_id` variable uniquely identifies all decisions. It indicates the type of the decision, the lead proceeding associated with the decision, and the date of the decision.
- You should always use the `decision_id` variable to merge tables. If you use `ecli` or `cellex` you may inadvertently lose or duplicate observations, as neither of those ID numbers is actually uniquely identifying. If you need to use them to merge in data, carefully check for duplicates or missed merges.
- The Court uses ECLI numbers to uniquely identify decisions, but ECLI numbers do not always uniquely identify decisions. There are two problems: (1) some ECLI numbers are assigned to more than one decision, and (2) some decisions have been assigned more than one ECLI number. Both problems particularly affect orders. Always use `decision_id` as the unique identifier.
- `lead_proceeding_id` refers to the lead proceeding, not all proceedings. When the Court joins proceedings together, a decision may apply to multiple proceedings. The `lead_proceeding_id` records only the lead proceeding, which is the one that was lodged first. To find all proceedings a decision applies to, you need to use the `joins` table in combination with `lead_proceeding_id`.
- A small number of orders in the Registry are excluded from the `decisions` table because the Registry does not report the date the order was published. Without a date, the decision cannot be uniquely identified. These orders also do not appear in InfoCuria or EUR-Lex, and their text is not available.

- The `decisions` table is coded based on four sources: The Registry, InfoCuria metadata, EUR-Lex metadata, and the text of decisions. These sources sometimes conflict (e.g., incorrect dates, incorrect ECLI numbers, incorrect document types, etc.). The data has been extensively cross-validated and corrected where necessary.

Variables

- `decision_id` `text` `primary key` `required`
A unique identifier for the decision. Format: `DOC_X_Y_####_####_Z_YYYYMMDD`, where `X` indicates the document type (`JUDG` for judgment, `ORD` for orders, and `OPIN` for AG opinions), `Y` indicates the court (`CJ` for the Court of Justice, `GC` for the General Court, and `CST` for the Civil Service Tribunal), the first number is the year of the case, the second number is the case number (with leading zeros), `Z` is an optional code for the proceeding suffix, and the final number is the date in the format `YYYYMMDD`. When two decisions of the same type are issued on the same day in the same case, the ECLI document number is appended in brackets to disambiguate (e.g., `_###`).
- `decision_date` `date` `required`
The date the decision was published. Format: `YYYY-MM-DD`.
- `decision_type` `text` `required`
The type of decision. Possible values: `AG opinion`, `AG view`, `decision`, `judgment`, `opinion`, `order`, and `ruling`.
- `court` `text` `required`
The court that issued the decision. Possible values: `Court of Justice`, `General Court`, and `Civil Service Tribunal`.
- `lead_proceeding_id` `text` `required`
The proceeding ID of the lead proceeding associated with this decision. When the Court has joined multiple proceedings together, this is the ID of the proceeding with the lowest proceeding number. To find all proceedings a decision applies to, use this field in combination with the `joins` table.
- `ecli` `text` `required`
The ECLI number for the decision. Format: `ECLI:EU:X:####:####`, where `X` indicates the court, the first number is the year, and the second number is a document number assigned by the Court (not predictable and not reconstructable from other metadata). ECLI numbers do not always uniquely identify decisions. Some decisions share an ECLI number; some decisions have multiple ECLI numbers. Always use `decision_id` for merging and filtering. Missing for decisions that do not appear in any official source with an ECLI number.

- `celex` `text` `nullable`

The CELEX number for the decision. Format: `6####XY####`, where the first number is the year, `XY` is the CELEX descriptor, where `X` indicates the court and `Y` indicates the decision type, and the second number is the case number with leading zeros. Some CELEX numbers include a suffix of the form `(##)` to differentiate multiple decisions of the same type in the same case. CELEX numbers do not always uniquely identify decisions. Missing for decisions that do not appear in EUR-Lex.

Procedures

One observation per outcome per legal procedure per decision. A single legal procedure can have multiple outcomes (e.g., both `unfounded` and `successful`).

Notes

- A single decision can involve more than one legal procedure (e.g., an action for annulment and an action for damages in the same case).
- A single procedure within a decision can also have more than one outcome (e.g., an application that is partially inadmissible and partially unfounded).
- Legal procedures include the major procedures outlined in the Treaty on the Functioning of the European Union (TFEU) — references for a preliminary ruling, actions for annulment, actions for failure to act, actions for failure to fulfill obligations, actions for damages, appeals, and staff cases — as well as a range of minor procedures such as applications for interim measures, applications for legal aid, and applications as to costs.
- This table has one row per outcome per procedure, not one row per decision. A decision with two procedures each having two outcomes will produce four rows. Before joining this table to any other decision-level table, aggregate it to the decision level if you want to avoid row multiplication.
- Two rows with the same `decision_id` and `procedure` but different `outcome` values are both valid. This indicates that the Court rendered a mixed decision with respect to that procedure.
- Decisions in proceedings concerning a reference for a preliminary ruling or an arbitration clause do not have a disposition. For these procedures, the `outcome` variable will be purposefully missing.
- The data on procedures and outcomes is based primarily on EUR-Lex metadata. For judgments, when EUR-Lex metadata is not available, the data is hand-coded based on the French-language text of the judgment. Not all orders have data.

Variables

- `decision_id` `text` `required`

The decision associated with this procedure and outcome.

- `procedure` `text` `required`

The legal procedure (e.g., 'reference for a preliminary ruling', 'action for annulment', 'appeal').

- `outcome` `text` `nullable`

The outcome of the procedure. Note that `agreement compatible` and `agreement incompatible` only apply to opinions of the Court (which are different than AG opinions). Missing when EUR-Lex metadata does not indicate an outcome with respect to the procedure. Possible values: `unfounded`, `successful`, `inadmissible`, `unnecessary`, `settled`, `agreement compatible`, and `agreement incompatible`.

Referrals

One observation per referral from a national court per decision (in decisions concerning a reference for a preliminary ruling).

Notes

- This table only applies to decisions and procedures that concern a reference for a preliminary ruling. If you are joining this table to `decisions`, you should filter `decisions` to only include judgments and to the reference for preliminary ruling procedure (via the `procedures` table) first. Left-joining `referrals` directly to `decisions` without filtering will result in many unmatched rows.
- `national_court_name` is the name of the national court exactly as it appears in each specific decision, which may differ across decisions due to translation, abbreviation, or historical naming changes. Use `national_court_id` for consistent identification across decisions.
- Joined proceedings in preliminary ruling can have multiple referring courts. In such cases, there will be multiple rows with the same `decision_id`. You may want to aggregate before joining to decision-level data.
- `date_referred` and `date_received` are coded based on the French-language text of the decisions. The decision does not always include this information, especially in earlier years.

Variables

- `decision_id` `text` `required`
The decision associated with this referral.
- `index` `integer` `required`
A sequential identifier for each referral within a decision.
- `national_court_id` `text` `required`
A standardized identifier for the referring national court. This identifier is consistent across all decisions.
- `national_court_name` `text` `required`
The name of the national court as it appears in the text of the decision.

- `member_state` `text` `nullable`

The EU member state in which the referring court is located. Missing for some referring courts that are not located within a member state.

- `date_referred` `date` `nullable`

The date the national court made the referral. Format: `YYYY-MM-DD`. This is coded based on the text of the judgment. Missing when the judgment does not report the date.

- `date_received` `date` `nullable`

The date the referral was formally received by the Court. Format: `YYYY-MM-DD`. This is coded based on the text of the judgment. Missing when the judgment does not report the date.

Hearings

One observation per hearing per decision.

Notes

- There can be multiple hearings per proceeding or no hearings. If a proceeding does not have a hearing, it does not appear in the `hearings` table, so not all decisions in the `decisions` table will have an observation in the `hearings` table.
- You can use hearing dates and decision dates to track how long it took for the proceeding to progress through the judicial process.

Variables

- `decision_id` `text` `required`
The decision associated with the hearing.
- `hearing_date` `date` `required`
The date of the hearing. Format: `YYYY-MM-DD`.

Directory Codes

One observation per EUR-Lex directory code per decision. Directory codes classify decisions by subject matter using a hierarchical coding system.

Notes

- The `directory_codes` table records the case law directory codes that the Court assigns to each judgment. The data is collected from EUR-Lex. Directory codes are a hierarchical classification system used by the Court and the Commission to organize CJEU case law by subject matter. There are two versions of the directory: a pre-Lisbon directory and a post-Lisbon directory, reflecting the reorganization of EU law following the Treaty of Lisbon.
- Each code is organized into a hierarchy: a section (top level), a major code (second level), and a minor code (third level, where applicable). The table records all levels of the hierarchy as both codes and human-readable labels.
- A decision can have multiple directory codes or no directory codes. If EUR-Lex does not have metadata on the directory codes for a decision, that decision will not be included in the `directory_codes` table. All missingness in this table is due to missingness in EUR-Lex. Since directory codes are assigned by the Court, they cannot be hand-coded to fill in missing data.
- Multiple rows per decision are common. A decision that touches on multiple areas of law will have multiple rows. You may want to aggregate before joining to decision-level data to avoid row multiplication.
- Both the pre-Lisbon and post-Lisbon case law directories can have many nested sub-categories. The `directory_codes` table codes truncates the codes after three levels, keeping a section, a major code (if available), and a minor code (if available).
- The pre-Lisbon and post-Lisbon case law directories are not directly comparable. The reorganization of the directory after the Treaty of Lisbon (2009) means that codes from before and after this date do not map cleanly to each other. When conducting longitudinal analyses, be careful to account for this break. Some codes can be easily mapped across, and some cannot.

Variables

- `decision_id` `text` `required`

The decision associated with this directory code.

- `code` `text` `required`
The full EUR-Lex directory code.
- `section_code` `text` `required`
The top-level section code.
- `section_label` `text` `required`
The label for the top-level section.
- `major_code` `text` `nullable`
The major subdivision code within the section. Missing when the directory code does not include a major code.
- `major_label` `text` `nullable`
The label for the major subdivision. Missing when the directory code does not include a major code.
- `minor_code` `text` `nullable`
The minor subdivision code within the major subdivision. Missing when the directory code does not include a minor code.
- `minor_label` `text` `nullable`
The label for the minor subdivision. Missing when the directory code does not include a minor code.

Keywords

One observation per EUR-Lex keyword per decision.

Notes

- The `keywords` table records the keywords assigned to each decision by EUR-Lex. Keywords are short subject-matter descriptors that classify the legal topics addressed in a decision. A single decision can have multiple keywords.
- A decision can have multiple keywords or no keywords. If EUR-Lex does not have metadata on the keywords for a decision, that decision will not be included in the `keywords` table. All missingness in this table is due to missingness in EUR-Lex. Since keywords are assigned by EUR-Lex, they cannot be hand-coded to fill in missing data.

Variables

- `decision_id` `text` `required`

The decision associated with this keyword.

- `keyword` `text` `required`

The EUR-Lex keyword, exactly as it appears on EUR-Lex.

Headnotes

One observation per headnote per decision.

Notes

- The `headnotes` table records the headnotes assigned to each decision by the Court. The data is collected from EUR-Lex. Headnotes are brief phrases that indicate the content of the decision.
- A decision can have multiple headnotes or no headnotes. If EUR-Lex does not have metadata on the headnotes for a decision, that decision will not be included in the `headnotes` table. All missingness in this table is due to missingness in EUR-Lex. Since headnotes are written by the Court, they cannot be hand-coded to fill in missing data.

Variables

- `decision_id` `text` `required`
The decision associated with this headnote.
- `headnote` `text` `required`
The text of the headnote, exactly as it appears on EUR-Lex.

Judges and Advocates-General

One observation per unique individual who has served as a judge and/or Advocate-General on one of the CJEU courts.

Notes

- The `judges_advocates_general` table includes data on every individual who has served as a judge or Advocate General (AG) on the Court of Justice, the General Court, or the Civil Service Tribunal. There is one observation per individual. The data is hand-coded based on the official biographies on the Court's website.
- The table records each individual's name (in multiple formats for display and matching purposes), member state, birth year, and gender. It serves as the master reference table for judicial personnel and is the primary key table for all judge- and AG-related joins throughout the database.
- Individuals who have served in multiple roles appear only once. A person who has served as both a judge and an AG (at the same or different courts) has only one row in this table. Role-specific information (court, position, dates) is in the `appointments` table.
- Use `actor_id` when joining this table to `appointments`, `judge_assignments`, `advocate_general_assignments`, or `judge_advocate_general_bios`. Do not match individuals by name, as name formatting varies across sources and decisions.
- Last names are not uniquely identifying (e.g., Higgins and Rossi both occur twice). The `display_label` and `display_label_ascii` variables are both uniquely identifying, in addition to `actor_id`.
- Some individuals' names include diacritics that are uncommon and many not render correctly in some environments. All names in the database use UTF-8 encoding. Use `display_label_ascii` if you are running into encoding problems.

Variables

- `actor_id` `text` `primary key` `required`
A unique identifier for the individual.

- `display_label` `text` `required`

A uniquely identifying display label. For most individuals, this is their last name. For the two sets of individuals who share a last name (`Rossi` and `Higgins`), this also includes their first initial (e.g., `Rossi, L.`).

- `display_label_ascii` `text` `required`

An ASCII version of `display_label` , with diacritical marks removed.

- `full_name` `text` `required`

The full name of the individual, including diacritical marks. This variable is uniquely identifying.

- `first_name` `text` `required`

The first name of the individual, including diacritical marks.

- `last_name` `text` `required`

The last name of the individual, including diacritical marks. This variable is not uniquely identifying.

- `member_state` `text` `required`

The EU member state that appointed the individual.

- `birth_year` `integer` `required`

The year of birth.

- `is_female` `boolean` `required`

Whether the individual is female.

Appointments

One observation per appointment per individual. Includes positions as judge, Advocate-General, President, or Vice-President. An individual can have multiple appointments (e.g., first as a judge, then as President).

Variables

- `actor_id` `text` `required`
The individual holding the appointment.
- `index` `integer` `required`
A sequential identifier for each appointment within an individual, assigned in chronological order.
- `court` `text` `required`
The court to which the individual was appointed. Possible values: `Court of Justice`, `General Court`, and `Civil Service Tribunal`.
- `position` `text` `required`
The position to which the individual was appointed. Possible values: `judge`, `Advocate-General`, `President`, and `Vice-President`.
- `is_current` `boolean` `required`
Whether this is a current appointment as of the most recent database update.
- `start_date` `date` `required`
The start date of the appointment. Format: `YYYY-MM-DD`.
- `end_date` `date` `nullable`
The end date of the appointment. Format: `YYYY-MM-DD`. Missing if the individual is currently serving.

Biographies

One observation per biographical entry per individual. Covers prior professional positions, memberships, degrees, and distinctions.

Notes

- The `biographies` table records structured biographical information for each judge and AG, hand-coded from their official biographies on the Court's website. There is one observation per biographical entry per individual. Each entry represents one item from the official biography, categorized by type: a professional position held, a membership (e.g., in a bar association or professional organization), an academic degree, or a distinction or honor received.
- For entries classified as professional positions, the table records the type of position (lawyer, civil servant, academic, politician, judge, or other) and the start and end years of the position. This structured biographical data enables researchers to study the professional backgrounds of CJEU members and how those backgrounds relate to judicial behavior.
- For all biographical entries, `start_year` and `end_year` can be missing. The official biographies often do not include precise years for all entries. Do not interpret missing years as indicating no experience, and do not assume that the duration of a position extends until the start of the next one that is listed.
- The coding scheme for `position_type` uses several categories: `lawyer` (working as a lawyer at a law firm, company, court, or governmental institution; includes law clerks and legal counsel at government institutions), `civil servant` (working in domestic or international governmental institutions in non-elected positions; includes positions at all ranks from government lawyers to ministers; excludes judges), `academic` (holding a permanent academic position; temporary or visiting positions are excluded), `politician` (holding an elected position at national or sub-national level, or serving as a Member of the European Parliament), `judge` (serving as a judge, including magistrate judges, at international, national, or sub-national level; CJEU service is included), and `other`.
- Some positions that involve work as a lawyer at a government institution. These are coded as `civil servant; lawyer`. This reflects cases where the primary role was civil service but the individual also performed legal functions.
- Entries where `entry_type = 'membership'` may include memberships in bar associations, professional societies, or academic bodies. These are not classified by `position_type`.

Variables

- `actor_id` `text` `required`

The individual this entry belongs to.

- `index` `integer` `required`

A sequential identifier for each entry within an individual.

- `entry_type` `text` `required`

The type of biographical entry. Possible values: `position`, `membership`, `degree`, and `distinction`.

- `description` `text` `required`

A plain text description of the entry. These descriptions are closely based on the language used in the official biographies, but are not verbatim quotes.

- `position_type` `text` `nullable`

The type of professional position. Missing when the `entry_type` is not `position`. Possible values: `lawyer`, `civil servant`, `academic`, `politician`, `judge`, `civil servant; lawyer`, and `other`.

- `start_year` `integer` `nullable`

The year the position or membership began, if applicable. Missing when not provided by the official bios.

- `end_year` `integer` `nullable`

The year the position or membership ended, if applicable. Missing when not provided by the official bios.

Judge Assignments

One observation per judge per decision. Records which judges sat on the panel for each decision and their role.

Notes

- The `judge_assignments` table records which judges sat on the panel that issued each decision and the positions of the judges, including which judge served as the Judge-Rapporteur. There is one observation per judge per decision. The table covers all judgments and relevant orders issued by the Court of Justice, the General Court, and the Civil Service Tribunal.
- Judge names and rapporteur status are coded based on the text of the judgment for most decisions. For early judgments, the data was hand-coded from the text. For later judgments, names are extracted using the French-language text. Judge-Rapporteur identification is cross-verified against EUR-Lex metadata, and discrepancies are resolved using the text.
- This table does not include assignments for decisions where the text is unavailable. A small number of decisions have no text available in any official source. These decisions will have no rows in `judge_assignments`.
- The identity of the Judge-Rapporteur is not always available. At the General Court, the Judge-Rapporteur is indicated consistently only from the mid-2000s onward. For earlier decisions, no judge on the panel will be marked as the judge-rapporteur. For the Court of Justice and Civil Service Tribunal, the Judge-Rapporteur is available for all decisions.
- Do not merge the `cases` table directly with the `decisions` table. Decisions are linked to proceedings, not directly to cases. To connect cases to decisions, join `cases` to `proceedings` to `decisions`. If you try to merge `cases` and `decisions` together using `case_id` and `proceeding_id`, you will drop all observations where the `proceeding_id` includes a suffix, which includes all appeals and some references for a preliminary ruling.
- The Judge-Rapporteur is verified to be on the panel. Panels are verified to have an odd number of judges. Panel sizes are checked for consistency against chamber assignment.

Variables

- `decision_id` `text` `required`

The decision the judge sat on.

- `judge_id` `text` `required`

The judge.

- `is_judge_rapporteur` `boolean` `required`

Whether the judge is the judge-rapporteur for the proceeding.

Advocate-General Assignments

One observation per Advocate-General per decision. Only applies to Court of Justice decisions where an AG was assigned.

Notes

- The `advocate_general_assignments` table records which Advocate General (AG) was assigned to each decision. There is one observation per AG per decision. Although rare, it is possible for more than one AG to be associated with a single decision (e.g., when an AG submits an opinion in joined proceedings that were originally handled by different AGs). For merging, do not assume that there is a one-to-one correspondence between AGs and decisions. Not all decisions have an AG assignment.
- AGs are assigned to proceedings, but this table is at the decision level. This is because the identity of the AG is indicated in the text of the decision. You can use the `decisions` table to associate AGs with proceedings.

Variables

- `decision_id` `text` `required`
The unique identifier for the decision.
- `advocate_general_id` `text` `required`
The unique identifier for the AG.

Chamber Assignments

One observation per chamber per decision. Records the formal chamber that heard the case.

Notes

- The `chamber_assignments` table records the chamber of the Court that the proceeding associated with the decision was assigned to. There is one observation per chamber per decision. This is a decision-level variable because the chamber assignment is indicated in the text of decisions. You can use the `decisions` table to associate the assigned chamber with the proceeding.
- A chamber is the standing configuration of the Court (e.g., the First Chamber, the Grand Chamber). This is different than a panel, which is the specific set of judges that actually participated in the proceeding and wrote the decision.
- Chamber names vary over time and across courts. The configuration and naming of chambers has changed throughout the Court's history. The same chamber name can refer to different sets of judges at different points in time. Longitudinal analyses using chamber data should account for these changes.
- The Grand Chamber and Plenary Assembly (the full Court) have quorum rules. The number of judges that sit on Grand Chamber and Full Court decisions varies, because these formations have minimum quorum rules that allow for some variation. The actual number of judges who participated in the proceeding is recorded in `judge_assignments`.

Variables

- `decision_id` `text` `required`
The unique identifier of the decision.
- `court` `text` `required`
The court of the chamber. Possible values: `Court of Justice`, `General Court`, and `Civil Service Tribunal`.
- `chamber` `text` `required`
The name of the chamber (e.g., 'Grand Chamber', 'Fifth Chamber'). See the `chambers` table for a full list.

Parties

One observation per party per decision.

Notes

- The `parties` table records the parties that appear in CJEU decisions. There is one observation per party per decision. The table covers judgments published by the Court of Justice, the General Court, and the Civil Service Tribunal that appear in the Registry, InfoCuria, or EUR-Lex. Party data is extracted from the text of judgments.
- For each party, the table records the party's name as it appears in the text of the specific document, a standardized party identifier, the party's type (e.g., EU member state, EU institution, firm, natural person), and the party's role in the proceeding (e.g., applicant, defendant, intervener).
- Most decisions involve multiple parties. You may want to aggregate before merging with decision-level data to avoid row multiplication, depending on your use case.
- Party names and party roles are coded based on the French-language text of the judgment, where available, using a carefully prompted LLM. Party type is hand-coded for EU-related actors (e.g., EU institutions, EU agencies, EU member states, etc.) and LLM-coded for other actors. It is possible for the LLM to make mistakes. LLM-coded data is hand-validated where possible.
- Interveners are included as parties. The `party_role` field distinguishes between applicants, defendants, interveners, and their counterparts in appeal proceedings (e.g., `applicant at first instance`, `intervener in the appeal`). If you want only the main parties, filter on `party_role`.
- The `parties` table also includes observers that are not technically parties in the proceeding. Their `role` is coded as `observer`. If a party also submits an observation, their main role in the case will take precedence in this table. See the `observers` table for data on all observers, including parties.
- `party_name` is the name as it appears in each specific decision. It is not standardized and can vary across documents due to language differences, abbreviation, or changes in institutional names over time. Do not use `party_name` for cross-decision matching. Use `party_id` instead.

- `party_id` is a standardized identifier that is consistent across all decisions in the database. It uniquely identifies each actor name, accounting for minor variations in spelling and punctuation. This helps to de-duplicate actors when there are slight variations in their names. However, there may still be actors in the database that have different names but are really the same actor. All EU-related actors have been completely de-duplicated (e.g., searching for the Commission using the Commission's unique ID will find all instances of the Commission, no matter how the name of the Commission appears in the text).
- Parties listed only in annexes are not included. In some decisions, parties (especially in large competition cases) are listed in an annex to the judgment rather than in the main text. These parties are not captured in this table.

Variables

- `decision_id` `text` `required`
The unique identifier for the decision.
- `index` `integer` `required`
A sequential index for each party within a decision.
- `party_id` `text` `required`
A standardized identifier for the actor, based on a slug of the actor's name (lower case, ASCII, hyphens for spaces). Use this for merging, not `party_name`.
- `party_name` `text` `required`
The name of the party exactly as it appears in the text of this decision. The name of an actor may differ across decisions. If you want to find all decisions that an actor appears in, use `party_id`, not `party_name`.
- `party_type` `text` `required`
The type of actor (e.g., company, EU institution, member state, individual). See the actor types type for a full list of possible values.
- `party_role` `text` `required`
The role of the party in the proceeding. Possible values: `applicant`, `applicant at first instance`, `appellant`, `defendant`, `defendant at first instance`, `intervener`, `intervener at first instance`, `intervener in the appeal`, and `observer`.

Observers

One observation per observer per decision. An observer is an actor that submits a written observation.

Notes

- The `observers` table records the observers that appear in CJEU decisions. There is one observer per party per decision. This table only applies to references for a preliminary ruling. Observer data is extracted from the text of judgments.
- If an observer is also a party in the case, the actor will also appear in the `parties` table.
- For each observer, the table records the observer's name as it appears in the text of the specific document, a standardized observer identifier and the observers's type (e.g., EU member state, EU institution, firm, natural person). Most observers are the Commission or EU member states.
- Observer names are coded based on the French-language text of the judgment, where available, using a carefully prompted LLM. Observer type is hand-coded for EU-related actors (e.g., EU institutions, EU agencies, EU member states, etc.) and LLM-coded for other actors. It is possible for the LLM to make mistakes. LLM-coded data is hand-validated where possible.
- `observer_name` is the name as it appears in each specific decision. It is not standardized and can vary across documents due to language differences, abbreviation, or changes in institutional names over time. Do not use `observer_name` for cross-decision matching. Use `observer_id` instead.
- `observer_id` is a standardized identifier that is consistent across all decisions in the database. It uniquely identifies each actor name, accounting for minor variations in spelling and punctuation. This helps to de-duplicate actors when there are slight variations in their names. However, there may still be actors in the database that have different names but are really the same actor. All EU-related actors have been completely de-duplicated (e.g., searching for the Commission using the Commission's unique ID will find all instances of the Commission, no matter how the name of the Commission appears in the text).

Variables

- `decision_id` `text` `required`

The decision that the observer appears in. Observers are only recorded in the text of decisions. They are connect to proceedings and cases via the decision.

- `index` `integer` `required`

A sequential index for each observer within a decision.

- `observer_id` `text` `required`

A standardized ID for the observer, based on a slug of the actor's name (lower case, ASCII, hyphens for spaces). Use this for merging, not `observer_name`.

- `observer_name` `text` `required`

The name of the observer as it appears in the text of the decision. May be in any EU language.

- `observer_type` `text` `required`

The type of observer (e.g., EU institution, member state, company).

Lawyers

One observation per lawyer per party or observer per decision.

Variables

- `decision_id` `text` `required`

The decision that the lawyer-client relationship appears in. Lawyers and their clients are only recorded in the text of decisions. They are connected to proceedings and cases via the decision.

- `index` `integer` `required`

A sequential index for each lawyer within the decision.

- `lawyer_id` `text` `required`

A standardized ID for the lawyer. Format: `ACT_LAW_[slug]`, where `slug` is a standardized version of the lawyer's name (lower case, ASCII, hyphens for spaces).

- `lawyer_name` `text` `required`

The name of the lawyer as it appears in the text of the decision.

- `client_id` `text` `required`

A standardized ID for the client (party or observer) that the lawyer represents in the proceeding, based on a slug of the actor's name (lower case, ASCII, hyphens for spaces).

- `client_name` `text` `required`

The name of the client as it appears in the text of the decision.

Citations

One observation per unique citation between judgments.

Notes

- This table records all citations between judgments. The `citing_decision_id` is the ID of the judgment that makes the citation, and `cited_decision_id` is the ID of the judgment being cited. The citations are extracted from the French-language text of each judgment by an algorithm.
- This table includes all direct citations that reference a case number. It does not include indirect citations (e.g, citations by case name only). However, for judgment-to-judgment citations, this almost never makes a difference because cited judgments are almost always formally cited (with a case number) at least once in the citing judgment.

Variables

- `citing_decision_id` `text` `required`
The judgment that contains the citation.
- `cited_decision_id` `text` `required`
The judgment being cited.